

CSC 421 - Artificial Intelligence

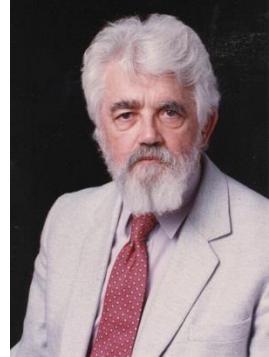
(Informal Introduction)

Alex Thomo

<http://webhome.cs.uvic.ca/~thomo>

Q&A

Father of AI:
John McCarthy
(September 4, 1927
– October 23, 2011)



Q. What is artificial intelligence?

A. Science and engineering of making intelligent machines.

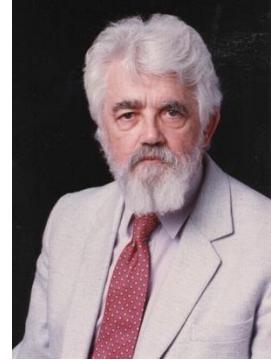
Q. Well...what exactly?

A. Computational part of achieving goals in the world.

Varying kinds and degrees of intelligence occur
in **people**,
many **animals**
some **machines**.

Q&A

Father of AI:
John McCarthy
(September 4, 1927
– October 23, 2011)



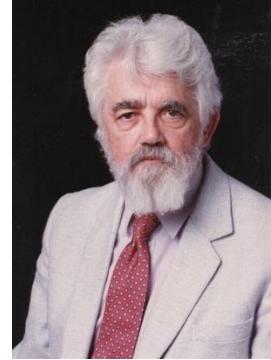
Q. When did AI research start?

A. After **WWII**, a people started to work on intelligent machines.

1. **Alan Turing** may have been the first (1947).
2. He suggested that AI was best researched by
programming computers rather than by building machines.

Q&A

Father of AI:
John McCarthy
(September 4, 1927
– October 23, 2011)



Q. What is the Turing test?

A. If the machine could successfully **pretend to be human**,
Not there yet!

Q. What about chess?

A. Chess programs now play at grandmaster level,
substituting large amounts of computation for understanding.

AI will be simple



"If it takes 200 years to achieve artificial intelligence and then finally there's a textbook that explains how it's done, the hardest part of that textbook to write will be the part that explains why people didn't think of it 200 years ago."

- John McCarthy

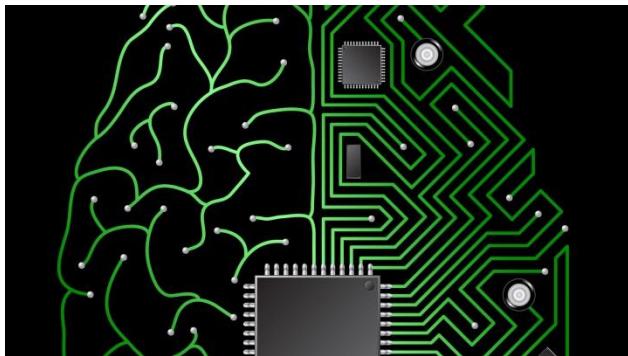
(Coined the term "Artificial Intelligence")

<https://www.youtube.com/watch?v=kL6J3y9ZCRQ>

What is AI

Build machines that

Think like people?



Think rationally?



Act like people?



Act rationally?



Some AI Topics

Topics I'll cover:

- Search
- Games
- Logic
- Reasoning under uncertainty
- Machine learning
- Computer Vision

Other topics:

- Robotics
- Natural Language Processing
- Etc.

Agents

- The job of AI is to design an **agent program** that implements the agent function--the mapping from percepts to actions.

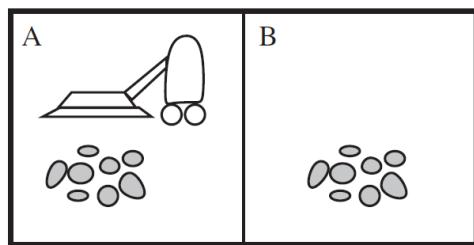
A very simple table-based agent

```
function TABLE-DRIVEN-AGENT(percept) returns an action
  persistent: percepts, a sequence, initially empty
              table, a table of actions, indexed by percept sequences, initially fully specified
  append percept to the end of percepts
  action  $\leftarrow$  LOOKUP(percepts, table)
  return action
```

- **Problem. Complexity:** Let P be the set of possible percepts and let T be the lifetime of the agent. The lookup table will contain $\sum_{t=1}^T |P|^t$ entries.

Reflex Agent

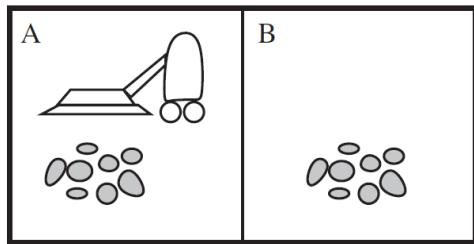
- Only consider the last percept.
- Solves the complexity problem, but can easily get stuck.



- Suppose a simple reflex vacuum agent has only a dirt sensor.
- Just two possible percepts: [Dirty] and [Clean].
- It can Suck in response to [Dirty];
- What should it do in response to [Clean]?
- Moving Left fails (forever) if it happens to start in square A.
- Moving Right fails (forever) if it happens to start in square B.

Model-based Reflex Agent

- Keep track of the current **state** of the world, using an internal model. Then choose an action in the same way as the reflex agent.



- Suppose a simple reflex vacuum agent has only a dirt sensor.
- Just two possible percepts: [Dirty] and [Clean].
- It can Suck in response to [Dirty]
- What should it do in response to [Clean]?
- If moving Left fails update state to “NoLeft”.
- If moving Right fails update state to “NoRight”.

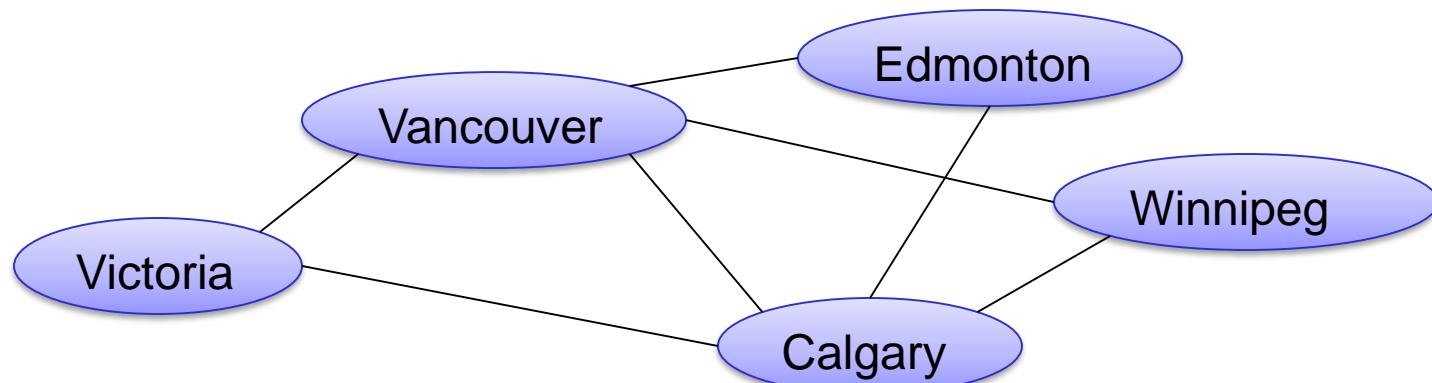
Search



Systematic **exploration of alternatives** for reaching a **goal**.
(Model-based agents with goal states)

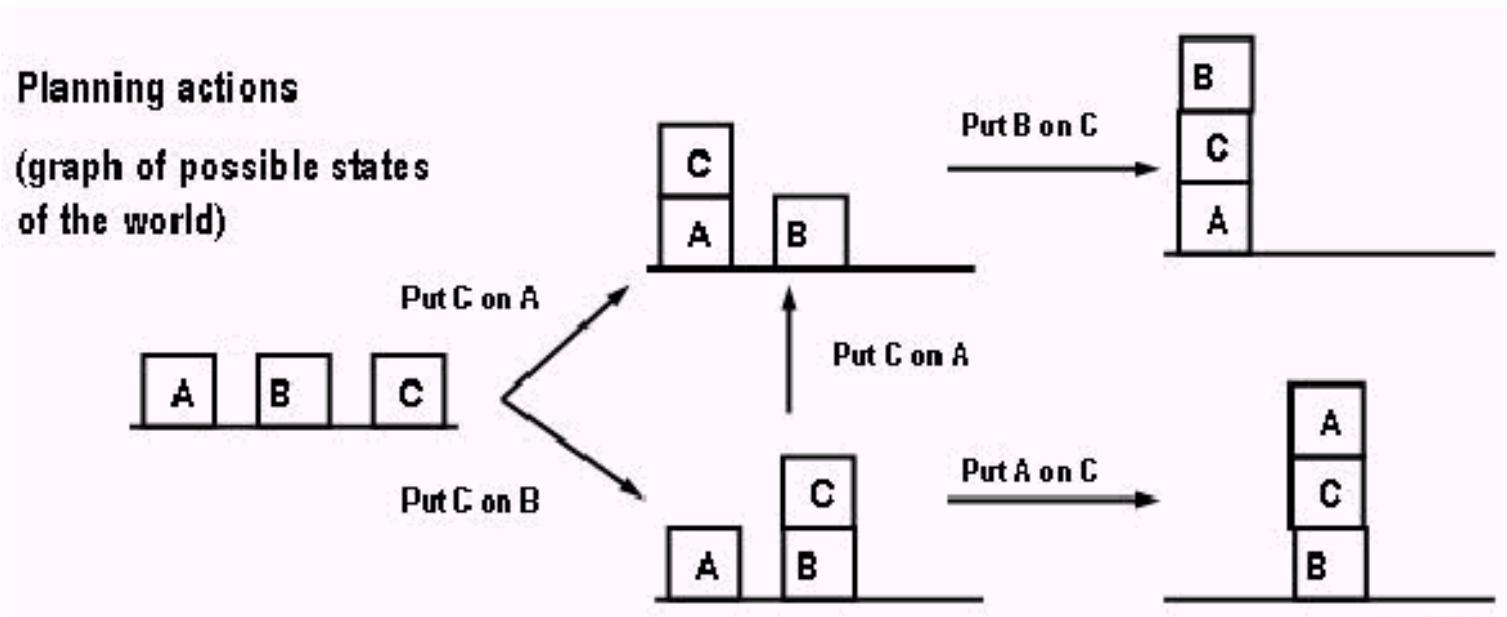
Search as path finding in graphs

Interested in **finding a path** through the graph that satisfies some property.



graph-node = state

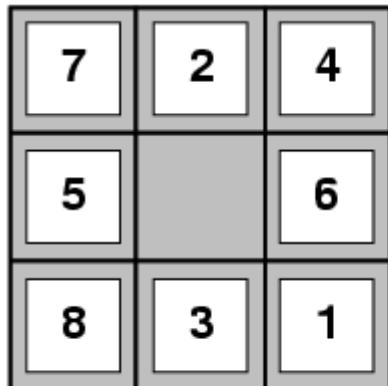
Graphs can be much more abstract



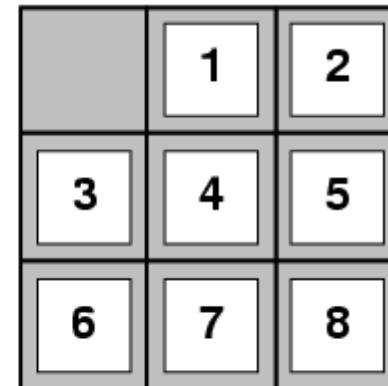
A **path** (from a start node to a goal node) is a "**plan of action**" to achieve a goal.

It's this type of graph that is of more general interest in AI.

Example: The 8-puzzle



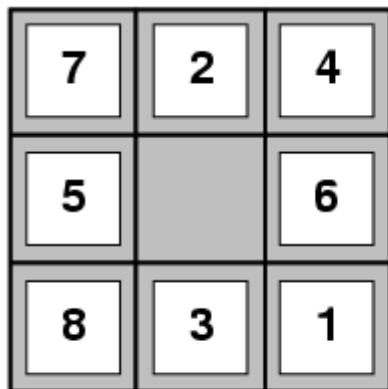
Start State



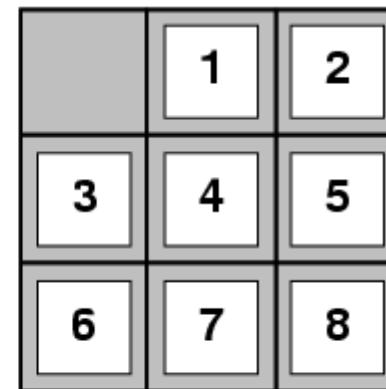
Goal State

- states?
- actions?
- goal test?
- path cost?

Example: The 8-puzzle



Start State

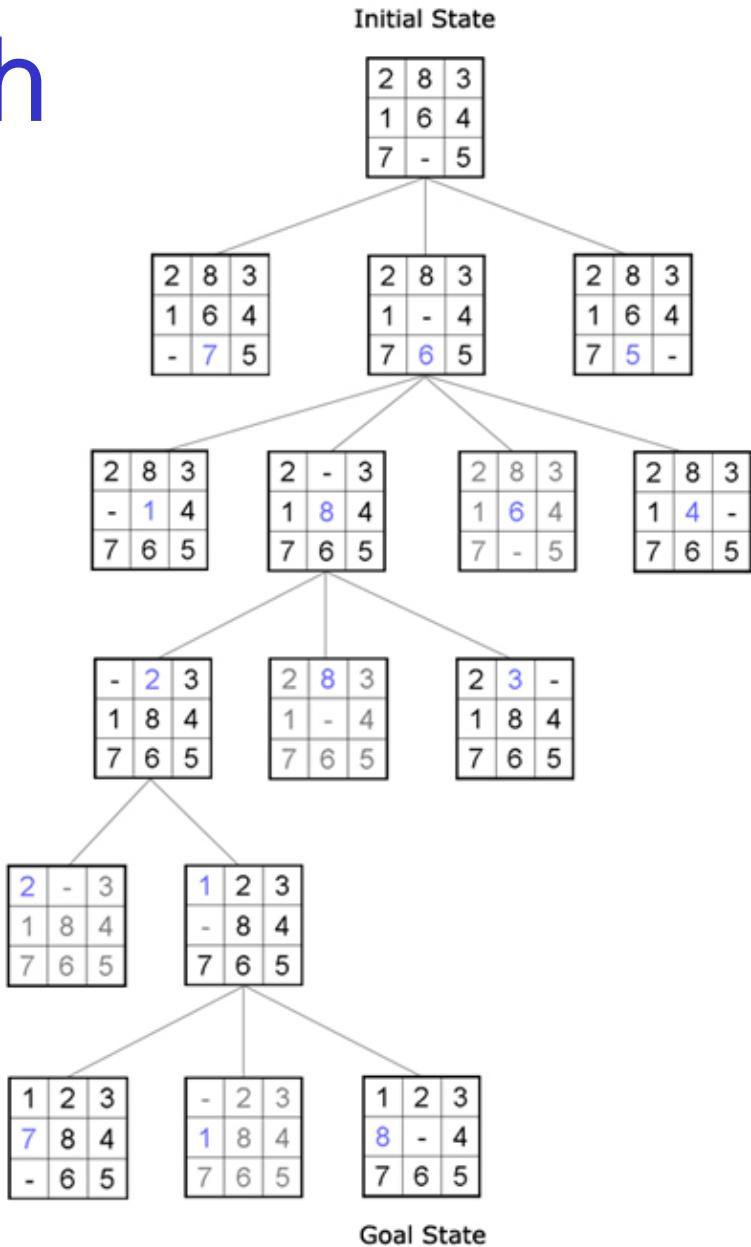


Goal State

- states? a board configuration is a state
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move

8-puzzle Search

State exploration



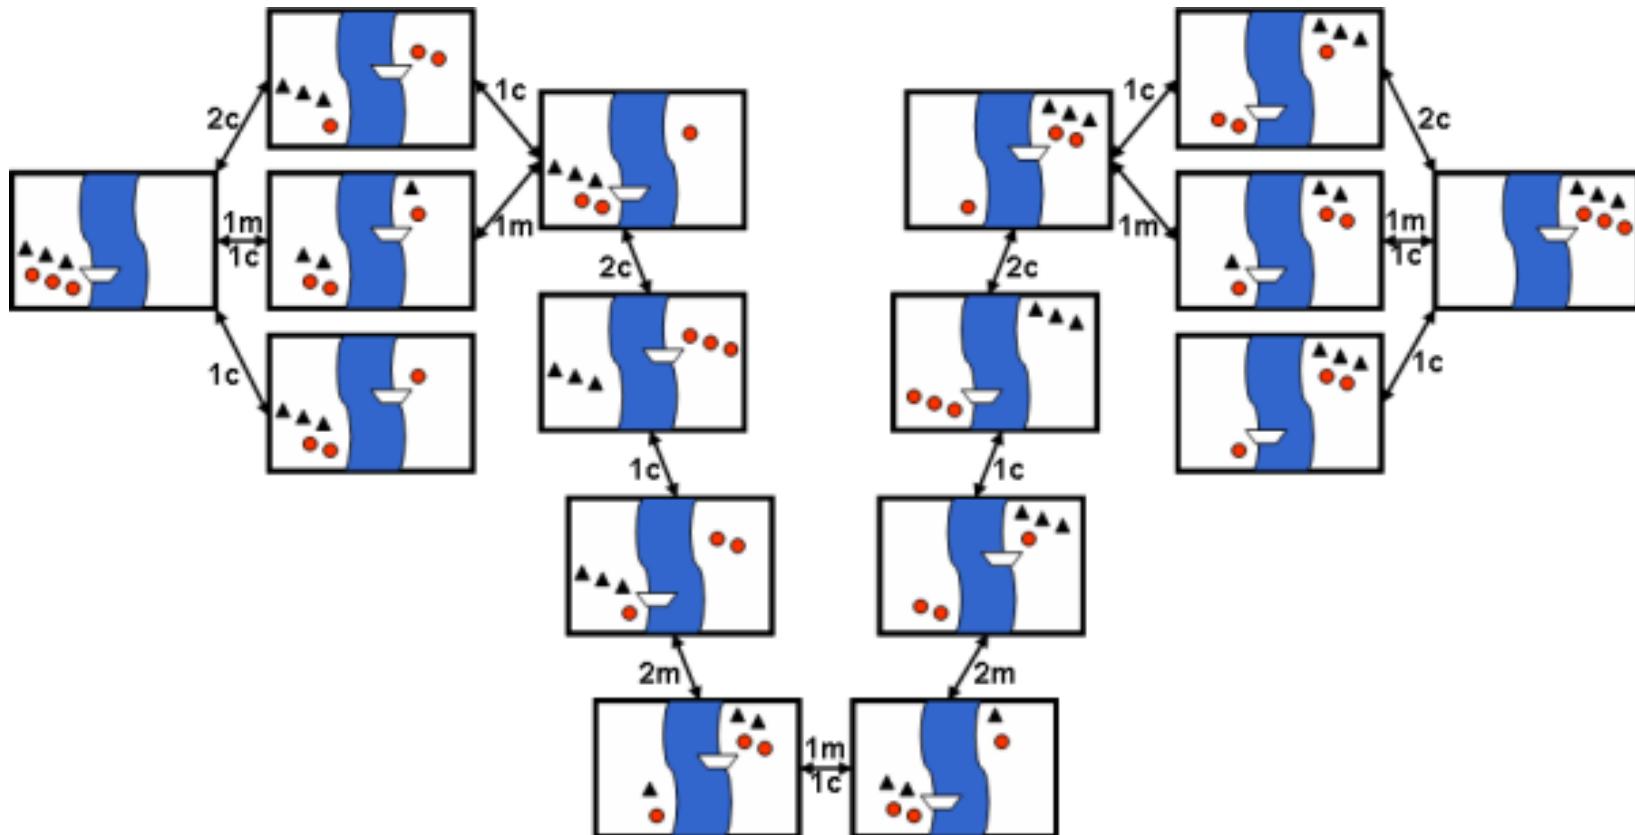
Missionaries and Cannibals



- Help 3 cannibals and 3 missionaries to move to the other side of the lake.
 - If there are more cannibals on one side than missionaries, the cannibals eat missionaries.

<http://www.learn4good.com/games/puzzle/boat.htm>

M&C Graph



M&C Solution (read down then left)

Move 2 cannibals to the left:



Move 2 missionaries to the left:



Move 2 cannibals to the left:



Move 1 cannibal back to the right:



Move 1 missionary and 1 cannibal back to the right:



Move 1 cannibal back to the right:



Move 2 cannibals to the left:



Move 2 missionaries to the left:



Move 2 cannibals to the left:



Move 1 cannibal back to the right:



Move 1 cannibal back to the right:



Logic

Human Logic

- Humans are, among other things, information processors.
- Our strength is the ability to represent and manipulate logical information.

A simple puzzle

Five blocks in a stack: **determine their exact arrangement.**

Premises

1. *The red block is on the green block.*
2. *The yellow block is on the green block or the blue block.*
3. *The green block is somewhere above the blue block.*
4. *There is some block on the black block.*

A simple puzzle

Five blocks in a stack: **determine their exact arrangement.**

Premises

1. *The red block is on the green block.*
2. *The yellow block is on the green block or the blue block.*
3. *The green block is somewhere above the blue block.*
4. *There is some block on the black block.*

Logical rules:

1. $(x \text{ on } z) \text{ and } (y \text{ on } z) \rightarrow x=y$
 $(x \neq y) \rightarrow \text{not}(x \text{ on } z) \text{ or } \text{not}(y \text{ on } z)$
2. $(x \text{ above } z) \rightarrow (x \text{ on } z) \text{ or } [(x \text{ above } y) \text{ and } (y \text{ on } z)]$
3. $(x \text{ on } y) \text{ and } (y \neq z) \rightarrow \text{not}(x \text{ on } z)$

Conclusions

1. *red on green.*
2. *yellow on blue.*
3. *green above yellow.*
4. *blue on black or green on black.*

Conclusions are proven using proofs with steps based on given facts, logical rules, and logical consequences.

Aristotle (384-322 B.C.E.)

What makes a step of a proof immediately obvious is its **form** rather than its **content**.



Examples

All Accords are Hondas.

All Hondas are Japanese.

Therefore, all Accords are Japanese.

All borogoves are slithy toves.

All slithy toves are flimsy.

Therefore, all borogoves are flimsy.

- In order to reach these conclusion, we don't need to know anything about *Hondas* and *Accords* or *borogoves* and *slithy toves* or what it means to be *flimsy*.
- What is interesting about these examples is that they share the same reasoning structure, with respect to the **pattern** shown below.

All x are y.

All y are z.

Therefore, all x are z.

Question

- Which patterns are sound?
- Well we just saw a sound pattern (previous slide)

Unsound Patterns

Pattern

All x are y.

Some y are z.

Therefore, some x are z.

Good Instance

All Toyotas are Japanese cars.

Some Japanese cars are made in America.

Therefore, some Toyotas are made in America.

Not-So-Good Instance

All Toyotas are cars.

Some cars are Porsches.

Therefore, some Toyotas are Porsches.

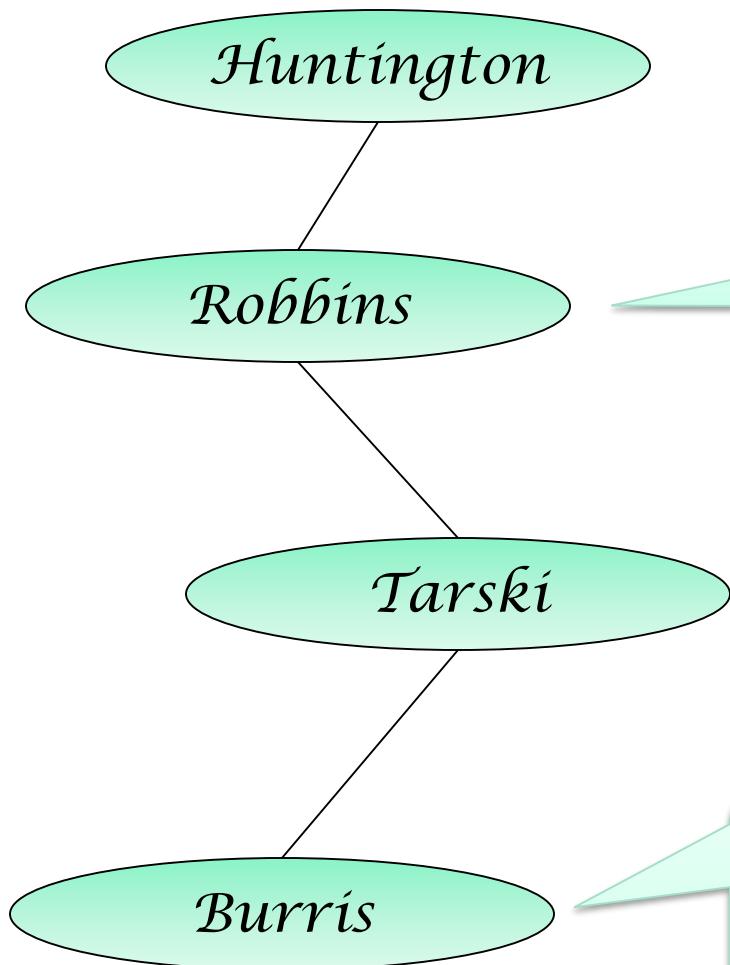
Formal Logic

1. Formal language for encoding information
2. **Legal** transformations

History

- **Theorem Provers** have come up with novel mathematical results.
 - The most famous of these concerns Robbins Algebra.
- In 1933, **E. V. Huntington** presented the following basis for Boolean algebra:
 - $x + y = y + x$. [commutativity]
 - $(x + y) + z = x + (y + z)$. [associativity]
 - $n(n(x) + y) + n(n(x) + n(y)) = x$. [Huntington equation]
- Shortly thereafter, **Robbins** conjectured that the Huntington equation can be replaced with a simpler one:
 - $n(n(x + y) + n(x + n(y))) = x$. [Robbins equation]

Gossip



Robbins said that he worked on the problem for some time, and then passed it on one of the century's most famous logicians, Dr. Albert Tarski of Stanford University.

Tarski, who is now dead, worked on the problem, included it in a book, and handed it out to graduate students and visitors.

Burris (University of Waterloo), for example, said that Tarski suggested the problem to him in the early 1970s, while he was visiting Stanford for a couple of months. Tarski, he said, "liked to throw out challenging problems to people passing through..."

October 10, 1996

- While mathematicians were batting around Robbins's problem, computer scientists were striving to see if they could get computers to reason.
- On October 10, 1996, after **eight days of computation**, EQP (a version of OTTER Theorem Prover) found a proof to Robbins conjecture.

Curiosity and the cat

- Everyone who loves all animals is loved by someone.
- Anyone who kills an animal is loved by no one.
- Jack loves all animals.
- Either Jack or Curiosity killed the cat, who is named Tuna.
- Did Curiosity kill the cat?

Curiosity and the cat in Prover 9

```
(animal(y) -> loves(x,y)) -> (exists y loves(y,x)).  
(exists y (animal(y) & kills(x,y)) -> (all z -loves(z,x))).  
animal(y) -> loves(jack,y).  
kills(jack,tuna) | kills(curiosity,tuna).  
cat(x) -> animal(x).  
cat(tuna).
```

Theorem:

```
kills(curiosity,tuna).
```

Prover9/Mace4

File Preferences View Help

Language Options Formulas Prover9 Options Mace4 Options Additional Input

Assumptions:

Highlight

Well Formed?

Clear

```
(animal(y) -> loves(x,y)) -> (exists y loves(y,x)).  
(exists y (animal(y) & kills(x,y)) -> (all z -loves(z,x))).  
animal(y) -> loves(jack,y).  
kills(jack,tuna) | kills(curiosity,tuna).  
cat(x) -> animal(x).  
cat(tuna).|
```

Show Current Input

Proof Search

Prover9

Time Limit: 60 seconds.

Start

Kill

State: Proof

Info

Show/Save

Model/Counterexample Search

Mace4

Time Limit: 60 seconds.

Start

Kill

State: Ready

Info

Show/Save

Nowadays

- AMD, Intel and others use **automated theorem proving** to verify that **optimized binary operations** correctly produce the intended result.



Uncertain

Representing Knowledge

Uncertain knowledge

- Typical example: **Diagnosis.**

Name	Toothache	...	Cavity
Smith	true	...	true
Mike	true	...	true
Mary	false	...	true
Quincy	true	...	false
...

- Can we certainly derive the **diagnostic rule**:
if $\text{Toothache}=\text{true}$ then $\text{Cavity}=\text{true}$

?

Isn't always right: Not all patients with toothache have cavities; some have gum disease, abscess, etc.

- We could try turning the rule into a **causal rule**:
 - if $\text{Cavity}=\text{true}$ then $\text{Toothache}=\text{true}$

Isn't always right: not all cavities cause pain.

Belief and Probability

- The connection between **toothaches** and **cavities** is not a logical consequence in either direction.
- However, we can provide a **degree of belief** on the rules. Our main tool for this is **probability theory**.

E.g. We believe there is, say, an 80% chance that the patient has cavity if he has a toothache.

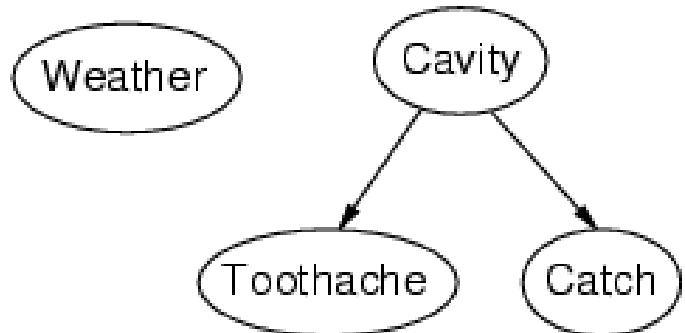
We usually get this belief from **statistical data or experience**.

Bayesian networks

A simple, **graphical** notation for conditional independence assertions.

Syntax:

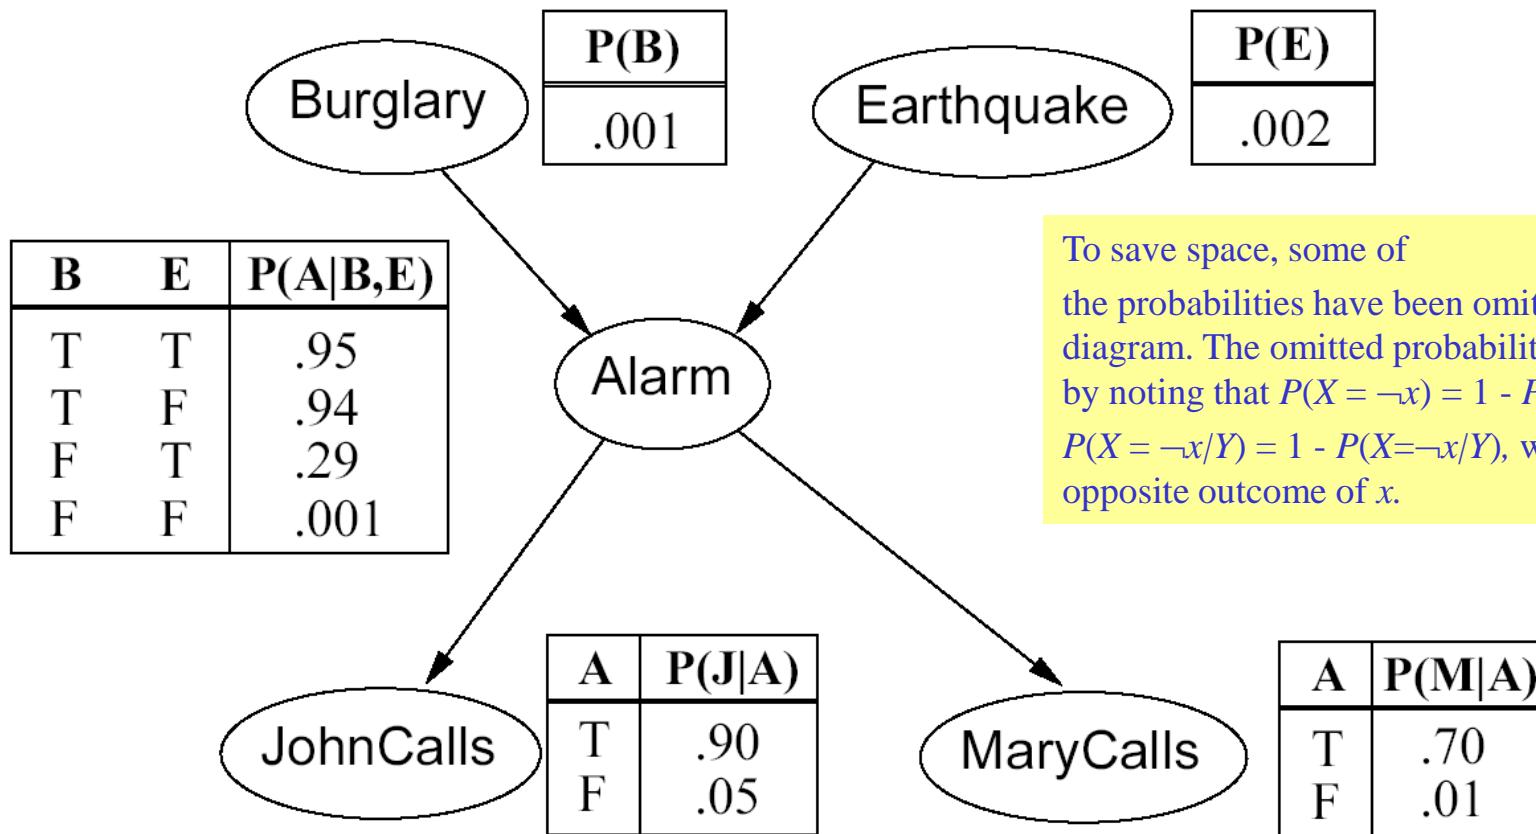
- a set of nodes, one per variable
- a directed, acyclic graph (**link means: "directly influences"**)
- **conditional probability tables** (CPTs) for each combination of parent values.



Example (Judea Perls' original example)

- I'm at work, neighbor John calls to say my alarm is ringing, but neighbor Mary doesn't call. Sometimes it's set off by minor earthquakes. **Is there a burglar?**
- John always calls when he hears the alarm, but sometimes confuses the telephone ringing with the alarm.
- Mary likes rather loud music and sometimes misses the alarm.
- Variables: *Burglary*, *Earthquake*, *Alarm*, *JohnCalls*, *MaryCalls*
- Network topology reflects knowledge:
 - A burglar can set the alarm off
 - An earthquake can set the alarm off
 - The alarm can cause Mary to call
 - The alarm can cause John to call

Example cont'd



To save space, some of the probabilities have been omitted from the diagram. The omitted probabilities can be recovered by noting that $P(X = \neg x) = 1 - P(X = x)$ and $P(X = \neg x|Y) = 1 - P(X=x|Y)$, where $\neg x$ denotes the opposite outcome of x .

The topology shows that burglary and earthquakes directly affect the probability of alarm, but whether Mary or John call depends only on the alarm.

Thus our assumptions are that they don't perceive any burglaries directly, and they don't confer before calling.

Inference

Example: $P(\text{burglary} \mid \text{johncalls}, \text{marycalls})$? (Abbr. $P(b|j,m)$)

$$P(b \mid j, m)$$

$$= \alpha P(b, j, m)$$

$$= \alpha \sum_a \sum_e P(b, j, m, a, e)$$

$$= \alpha (P(b, j, m, a, e) + P(b, j, m, \neg a, e) + P(b, j, m, a, \neg e) + P(b, j, m, \neg a, \neg e))$$

Computed as:

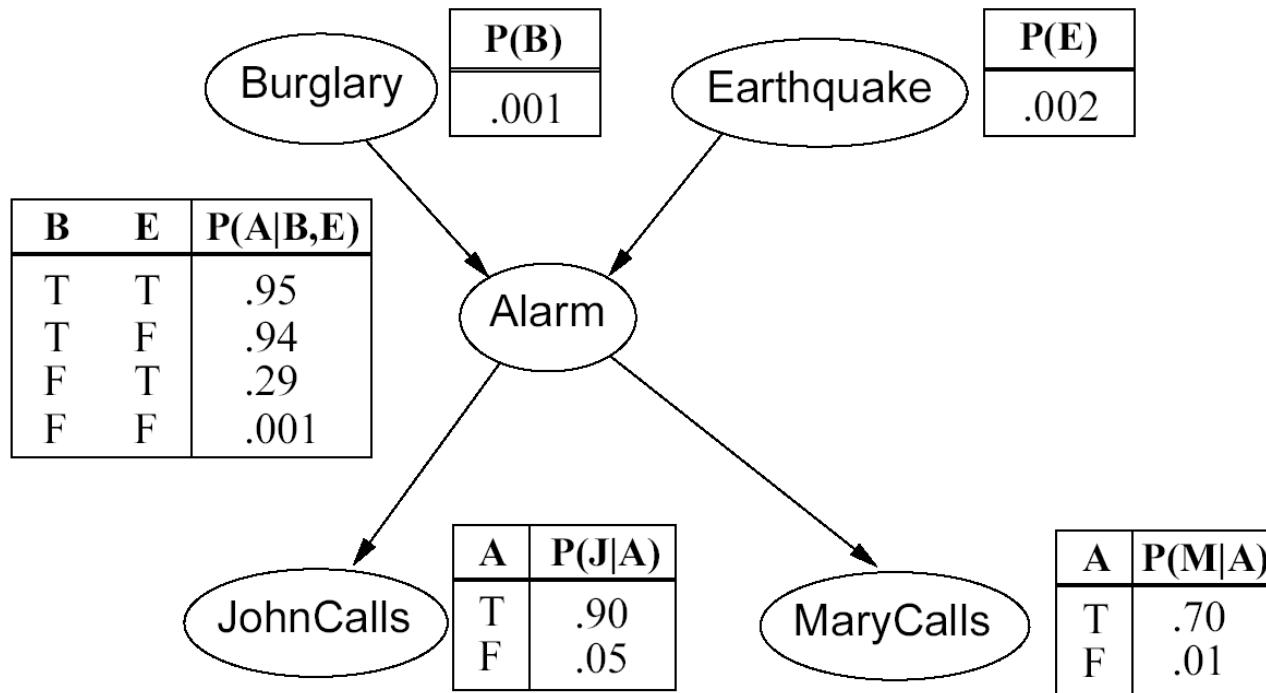
$$P(b)P(e)P(j|a)P(m|a)P(a|b, e)P(e)$$

Computed as:

$$P(b)P(e)P(j|\neg a)P(m|\neg a)P(\neg a|b, e)P(e)$$

Probabilities are taken from the tables.

Numerically...



$$P(b | j, m) = \dots = \alpha * 0.00059$$

$$P(\neg b | j, m) = \dots = \alpha * 0.0015$$

$$\mathbf{P(B | j, m)} = \alpha <0.00059, 0.0015> = \mathbf{<0.28, 0.72>}.$$

Aispace Demo

Belief and Decision Network Tool Version 5.1.10 --- untitled.xml

File Edit View Network Options Help

Make Observation Query P(e) Query Toggle Monitoring Select View Probability Table View/Modify Decision Add No-forgetting Arcs Optimize Decisions Independence Quiz

Create Solve

Click on a node to query its probability or utility.

```
graph LR; Burglary((Burglary)) --> Alarm((Alarm)); Earthquake((Earthquake)) --> Alarm; Alarm --> JohnCalls((JohnCalls  
Observed Value: T)); Alarm --> MaryCalls((MaryCalls  
Observed Value: T));
```

Query Results

Query Results for Variable Burglary [JohnCalls=T] [MaryCalls=T]

P (Burglary = T) = 0.28417

P (Burglary = F) = 0.71583

Belief Network Mode

M Microsoft PowerPoint...

OK

Nice Bayesian Network Example

<https://www.youtube.com/watch?v=4fcqyzVJwHM>

Right outside the window lay a dead Mr. Boddy. Alarmed, she called the police, and a detective was assigned to the case.

The detective, however, was not the most motivated man at the station, and always tried to make his job as easy as possible.



Learning

Learning

It's not always easy to instruct the machine what to do. Often it's best to let the machine "**learn**" how the world works.

We can think of at least three different problems being involved in learning:

- memory
- averaging
- generalization

Example problem

(Adapted from Leslie Kaelbling's example in the MIT courseware)

- Imagine I'm trying predict whether my neighbor is going to drive into work.
- Whether she drives into work seems to depend on the following attributes of the day:
 - temperature
 - expected precipitation
 - day of the week
 - what she's wearing

Memory

- Okay. Let's say we observe our neighbor on three days:

Temp	Precip	Day	Clothes	
25	None	Sat	Casual	Walk
-5	Snow	Mon	Casual	Drive
15	Snow	Mon	Casual	Walk

Memory

- Now, we find ourselves on a snowy “-5” degree Monday, and the neighbor is wearing casual clothes.
- Do you think she's going to drive?**



Temp	Precip	Day	Clothes	
25	None	Sat	Casual	Walk
-5	Snow	Mon	Casual	Drive
15	Snow	Mon	Casual	Walk
-5	Snow	Mon	Casual	

Memory

- Standard answer in this case is "yes".
 - This day is just like one of the ones we've seen before, and so it seems like a good bet to predict "yes."
- This is the most rudimentary form of learning, which is just to memorize the things you've seen before.



Temp	Precip	Day	Clothes	
25	None	Sat	Casual	Walk
-5	Snow	Mon	Casual	Drive
15	Snow	Mon	Casual	Walk
-5	Snow	Mon	Casual	Drive

Noisy Data

- Things aren't always as easy as they were in the previous case. What if you get this set of noisy data?

Temp	Precip	Day	Clothes	
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Drive
25	None	Sat	Casual	Drive
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Walk
25	None	Sat	Casual	?

- We have certainly seen this case before, but the problem is that it has had different answers. Our neighbor is not entirely reliable.

Averaging

- One strategy would be to predict the majority outcome.

Temp	Precip	Day	Clothes	
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Drive
25	None	Sat	Casual	Drive
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Walk
25	None	Sat	Casual	Walk

Generalization

- Dealing with previously unseen cases
- Will she walk or drive?

Temp	Precip	Day	Clothes	
22	None	Fri	Casual	Walk
3	None	Sun	Casual	Walk
10	Rain	Wed	Casual	Walk
30	None	Mon	Casual	Drive
20	None	Sat	Formal	Drive
25	None	Sat	Casual	Drive
-5	Snow	Mon	Casual	Drive
27	None	Tue	Casual	Drive
24	Rain	Mon	Casual	?

We might plausibly make any of the following arguments:

- She's going to walk because it's raining today and the only other time it rained, she walked.
- She's going to drive because she has always driven on Mondays...

Image Recognition

airplane



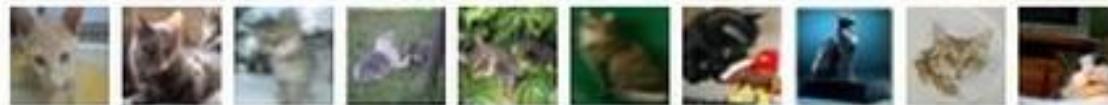
automobile



bird



cat



deer



dog



frog



horse

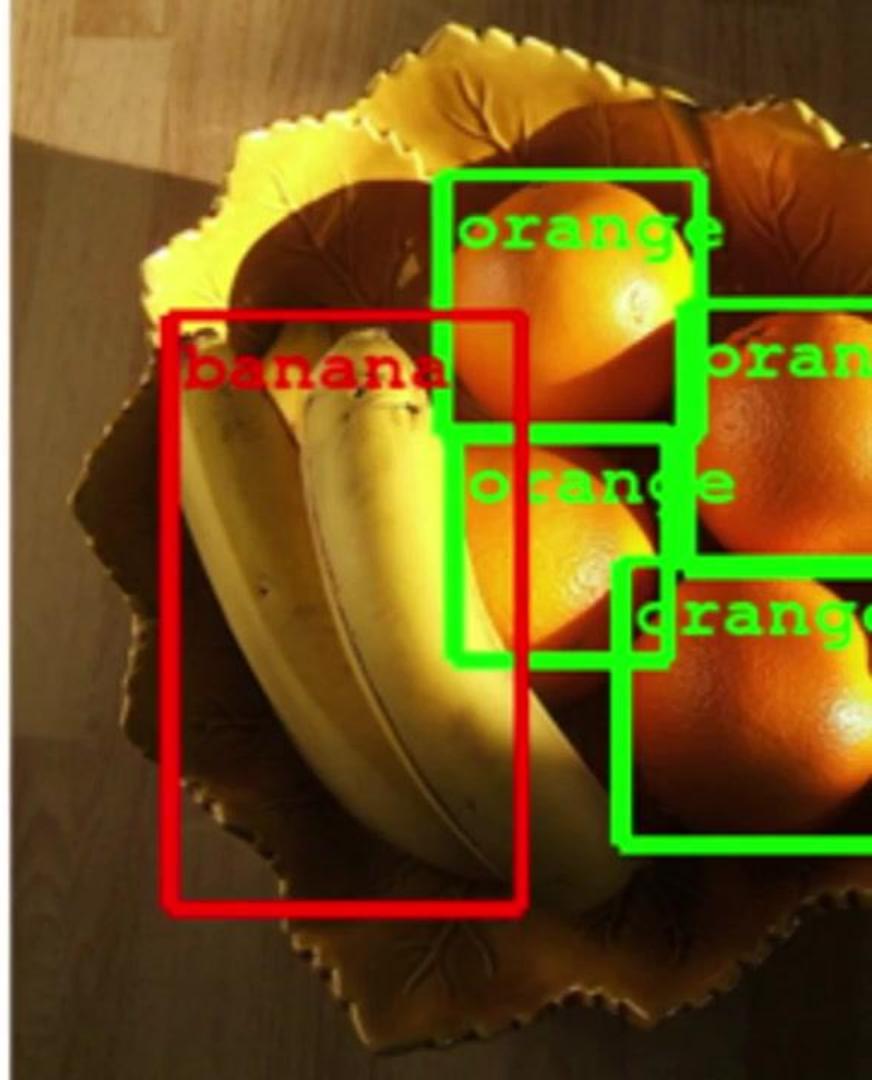
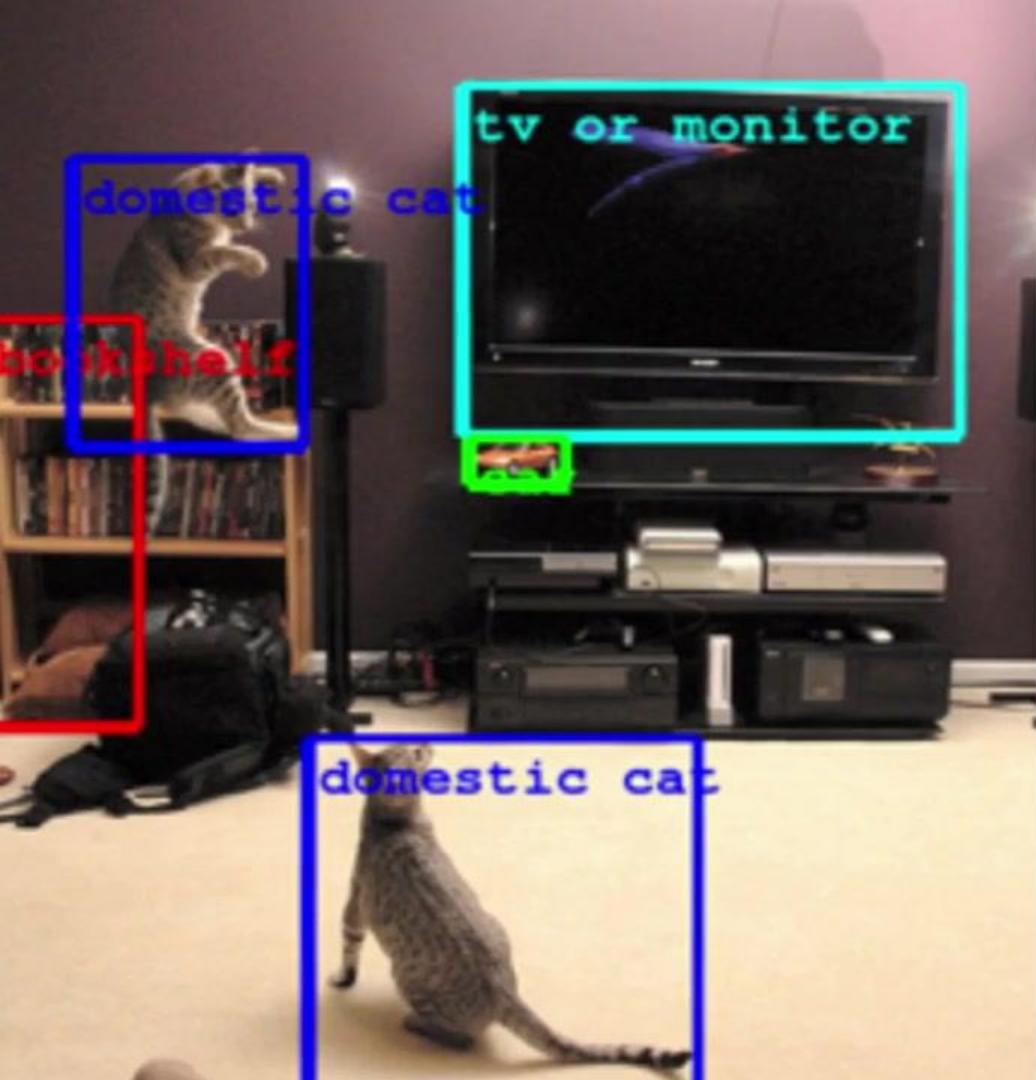


ship



truck





Taste of neural networks for OCR

For informal intro see:

<https://www.youtube.com/watch?v=ARODjRbGbSg>

The next slides are from this video.

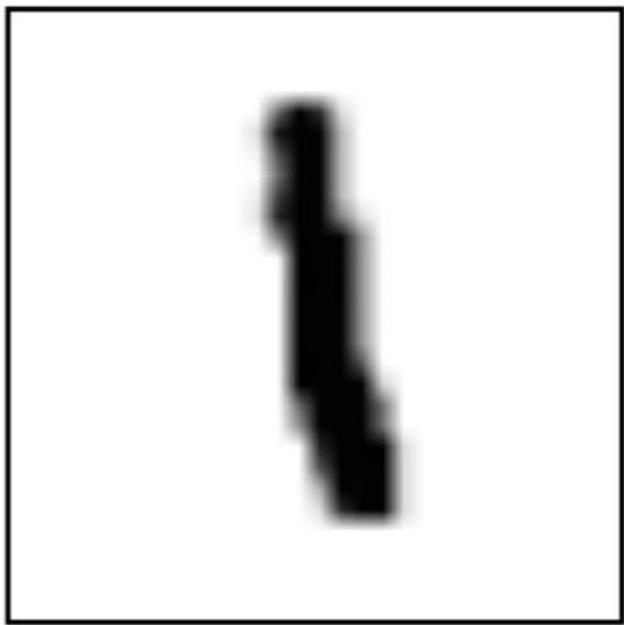
Few handwritten digits

504 / 92

70,000 of them

4377829931573776434337715892804718
51898976745274796044905143928085105
03643189262731004854777034572639488
3604341639093006328336847993775208
152817036924297119308146485776738791
77457447232507551269012752330042409
74053714190432263882135608771592290
52245374682930984293281391882500463
72209995144380831589574035051802078
963868686814924978955794805314152665
73123845338131421761254417350023431
12912424710467467691195319028458798
96706767939038251407028004453004630
25581188019257601040153642666276944
14143217884308679884810145150818574
48056138086724143479986976035116370
29800270575567965816172603263518470
97447257494992614235437115576062641
21602346049800830749873379842874592

One of them



~

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.6	.8	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.7	.1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.7	.1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.5	.1	.4	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.4	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.4	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.7	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.1	.1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	.9	.1	.1	0	0	0	0	0	0
0	0	0	0	0	0	0	.3	.1	.1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

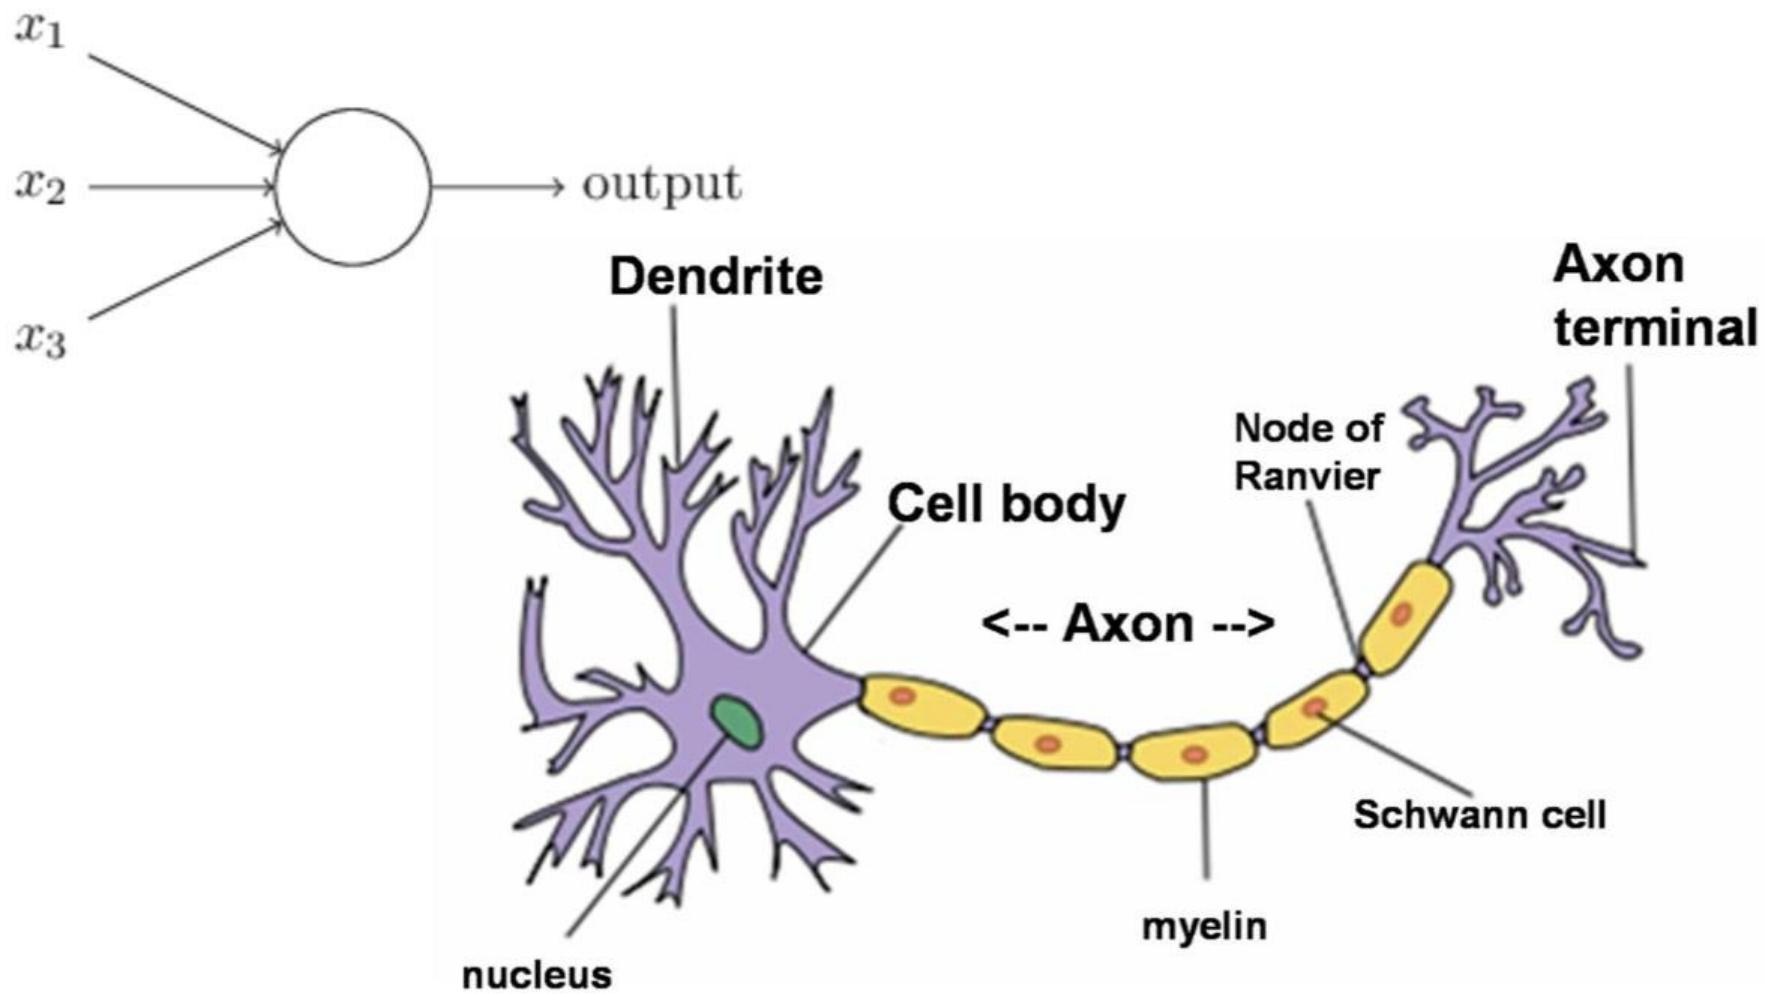
Not drawn to scale

The “classes”

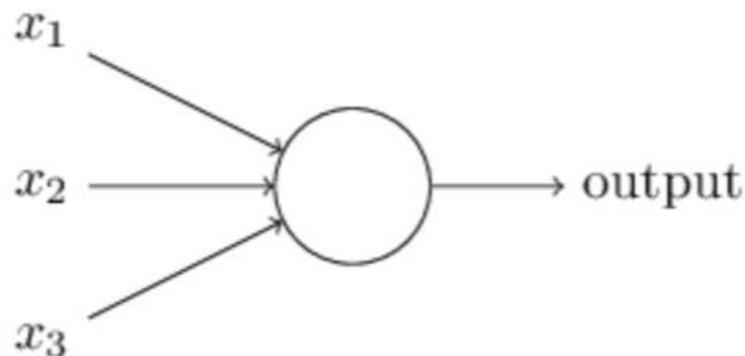
0 1 2 3 4

5 6 7 8 9

An artificial neuron and ... a natural one

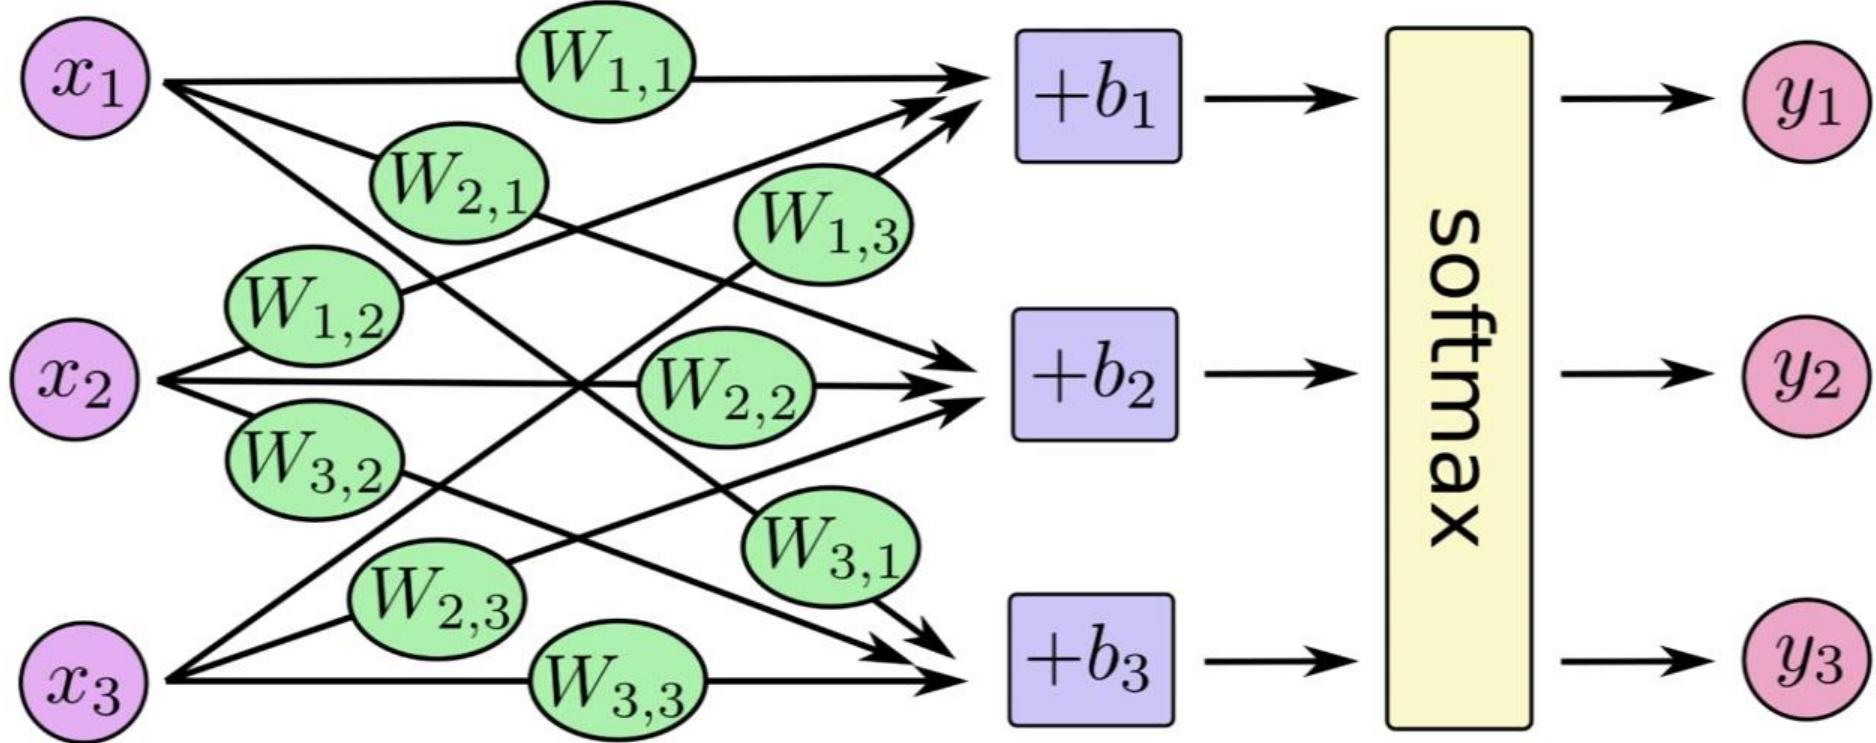


Mathematically speaking



$$\sum x_i W_i + b$$

A simple neural network



Inside a “class” neuron

$$evidence_c = \sum_{i=1}^{784} x_i W_{c,i} + b_c$$

Output vector

<0.01,0.91,0.01,0.01,0.01,0.01,0.01,0.01,0.01,0.01>

Output class –

the one corresponding to the highest value

<0.01,0.91,0.01,0.01,0.01,0.01,0.01,0.01,0.01,0.01>

0 1 2 3 4 5 6 7 8 9

Are these probabilities?

Obtaining probabilities –

normalizing the output vector

$$normevidence_c = \frac{evidence_c}{\sum_{c=0}^9 evidence_c}$$

A better normalization –

through amplification first

$$normevidence_c = \frac{e^{evidence_c}}{\sum_{c=0}^9 e^{evidence_c}}$$

A complex network graph is visible in the background, consisting of numerous small blue circular nodes connected by a web of thin blue lines. The graph is set against a dark, almost black, background.

Calculating the Weights

Compare probability distributions (vectors)

<0,1,0,0,0,0,0,0,0,0>

<0.01,0.91,0.01,0.01,0.01,0.01,0.01,0.01,0.01,0.01>

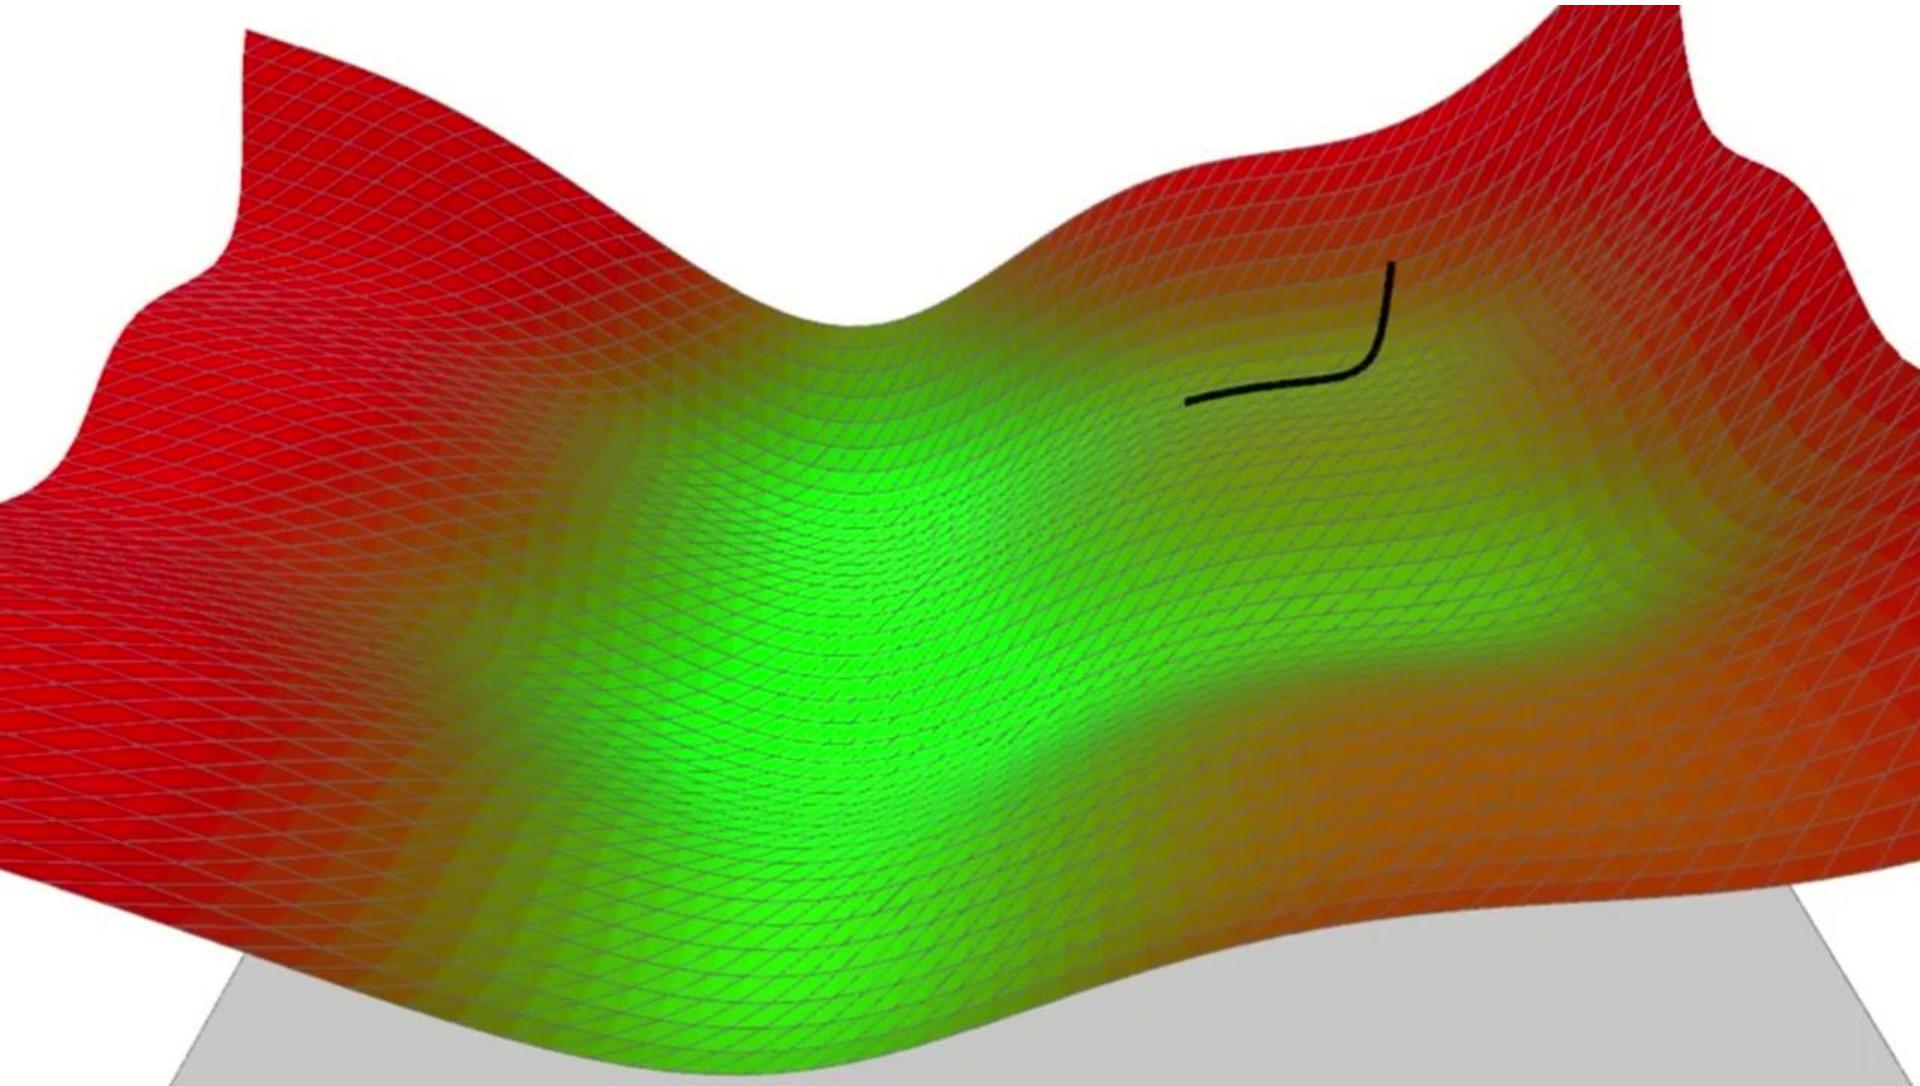
The Loss function

– log(probability)

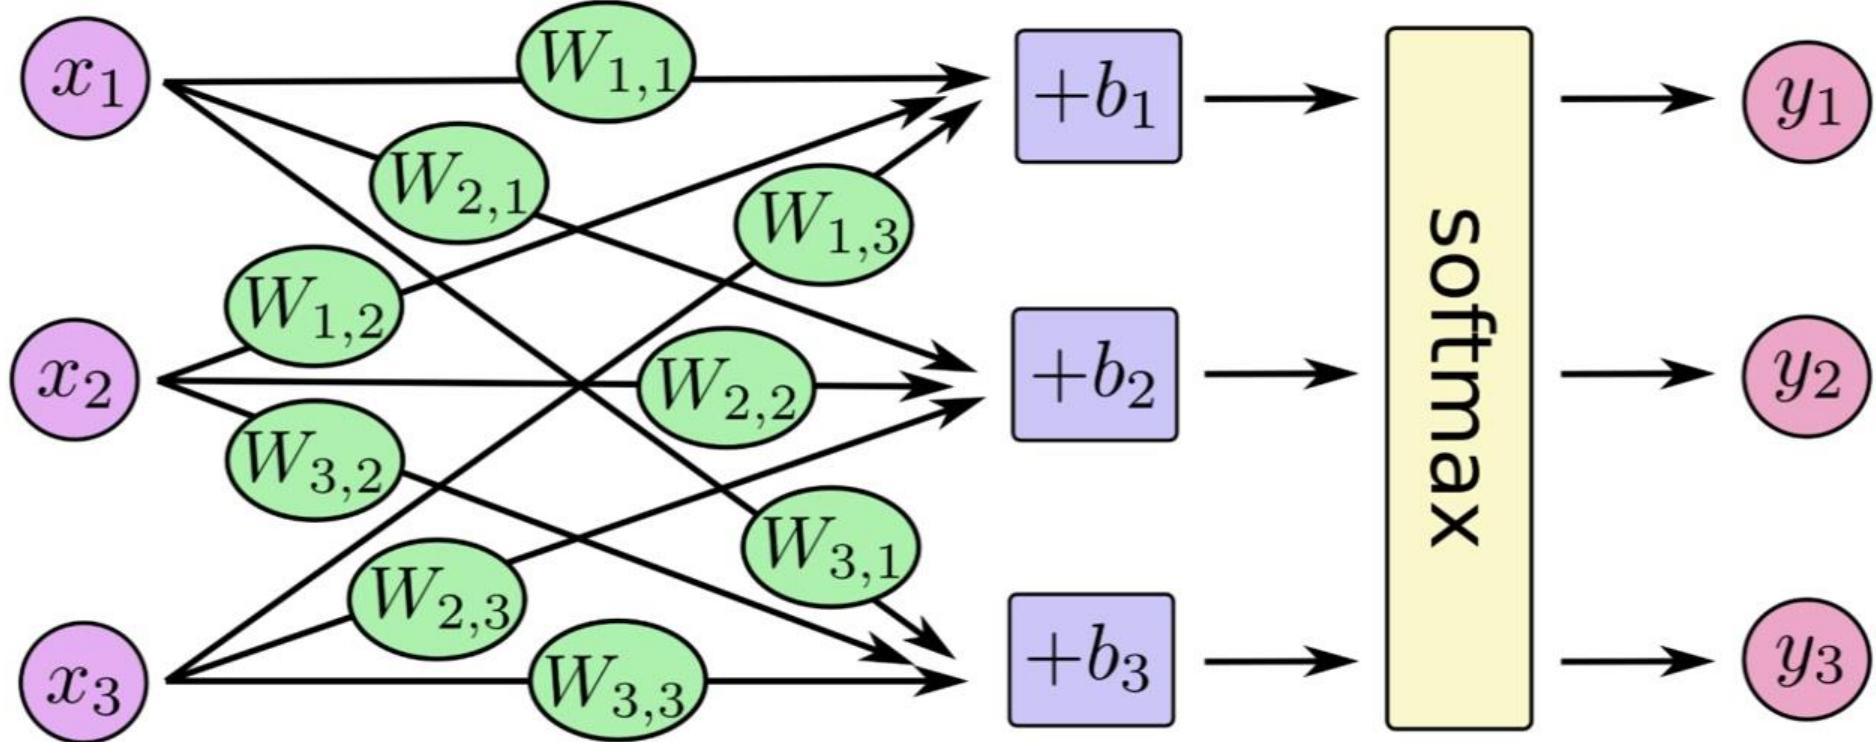
<0,1,0,0,0,0,0,0,0,0>

<0.01,0.91,0.01,0.01,0.01,0.01,0.01,0.01,0.01,0.01>

Navigating the Loss surface to reach the minimum

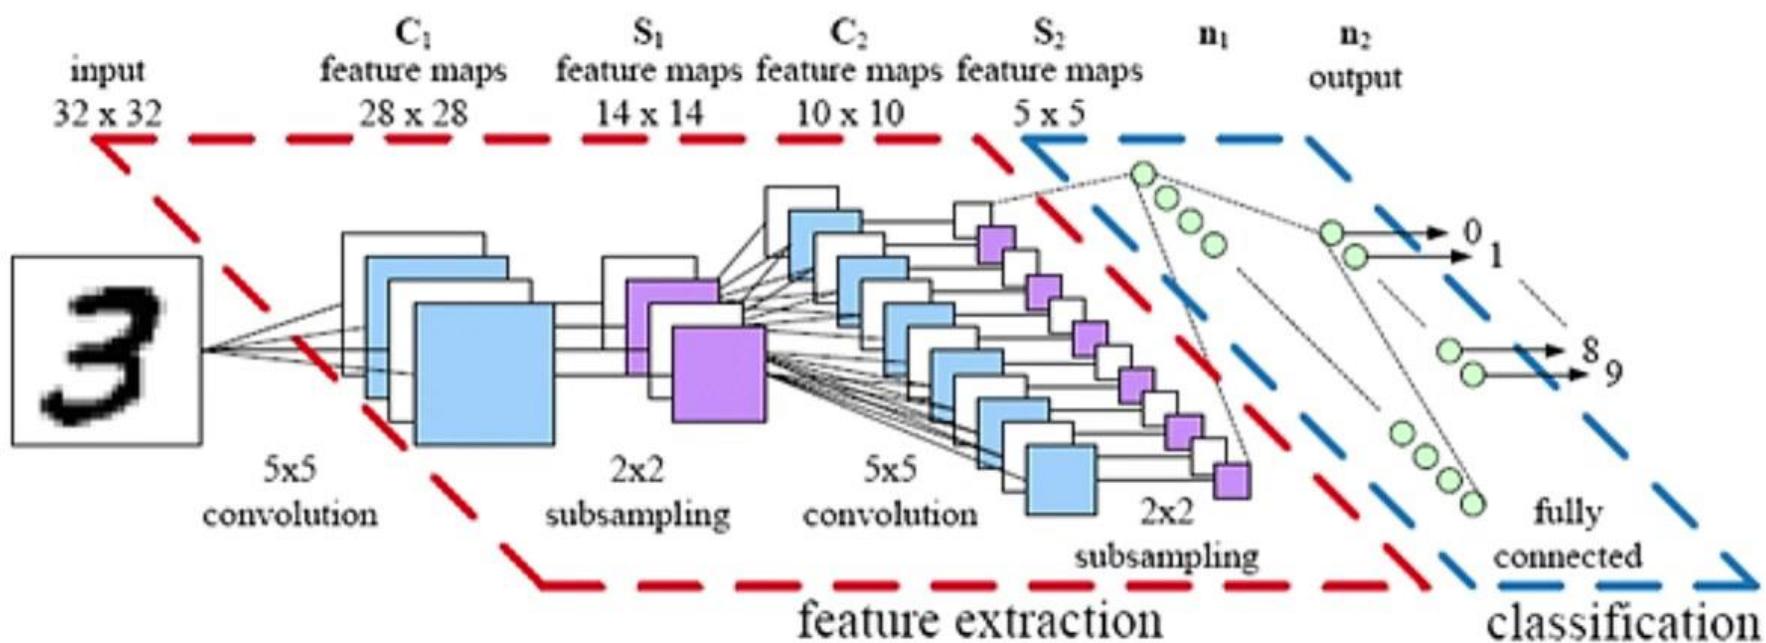


Accuracy of the simple NN



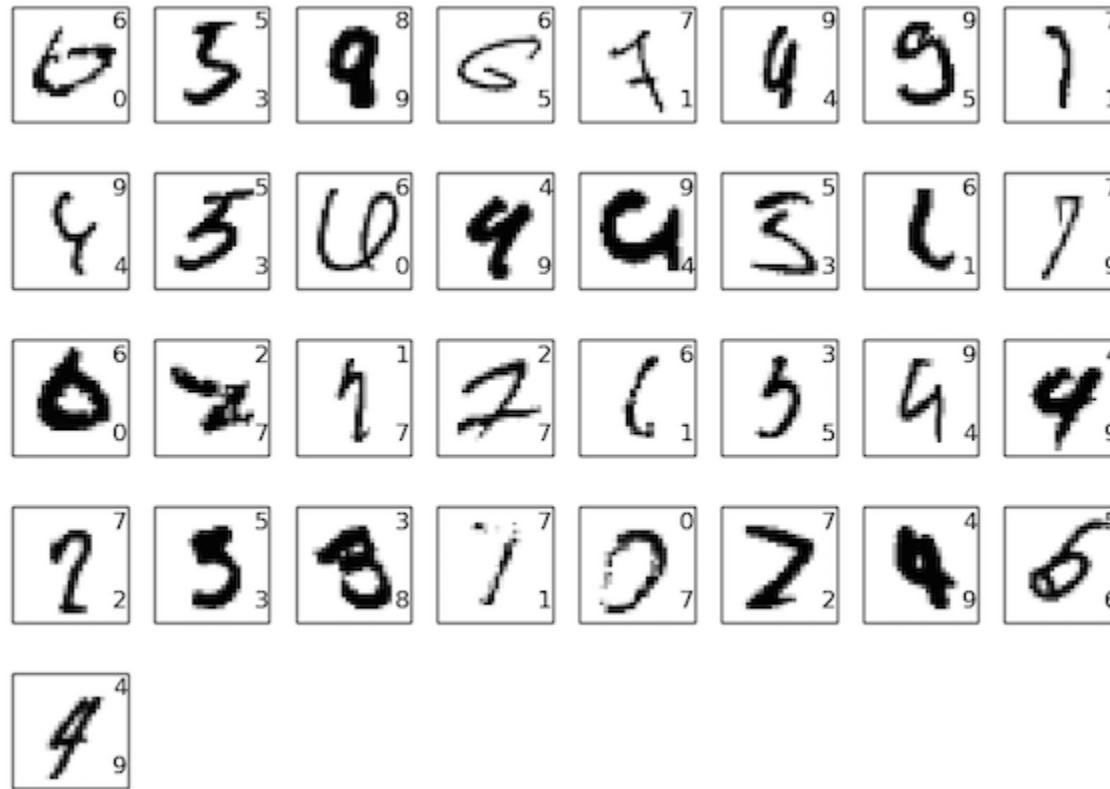
95%

More complicated (convoluted) NN



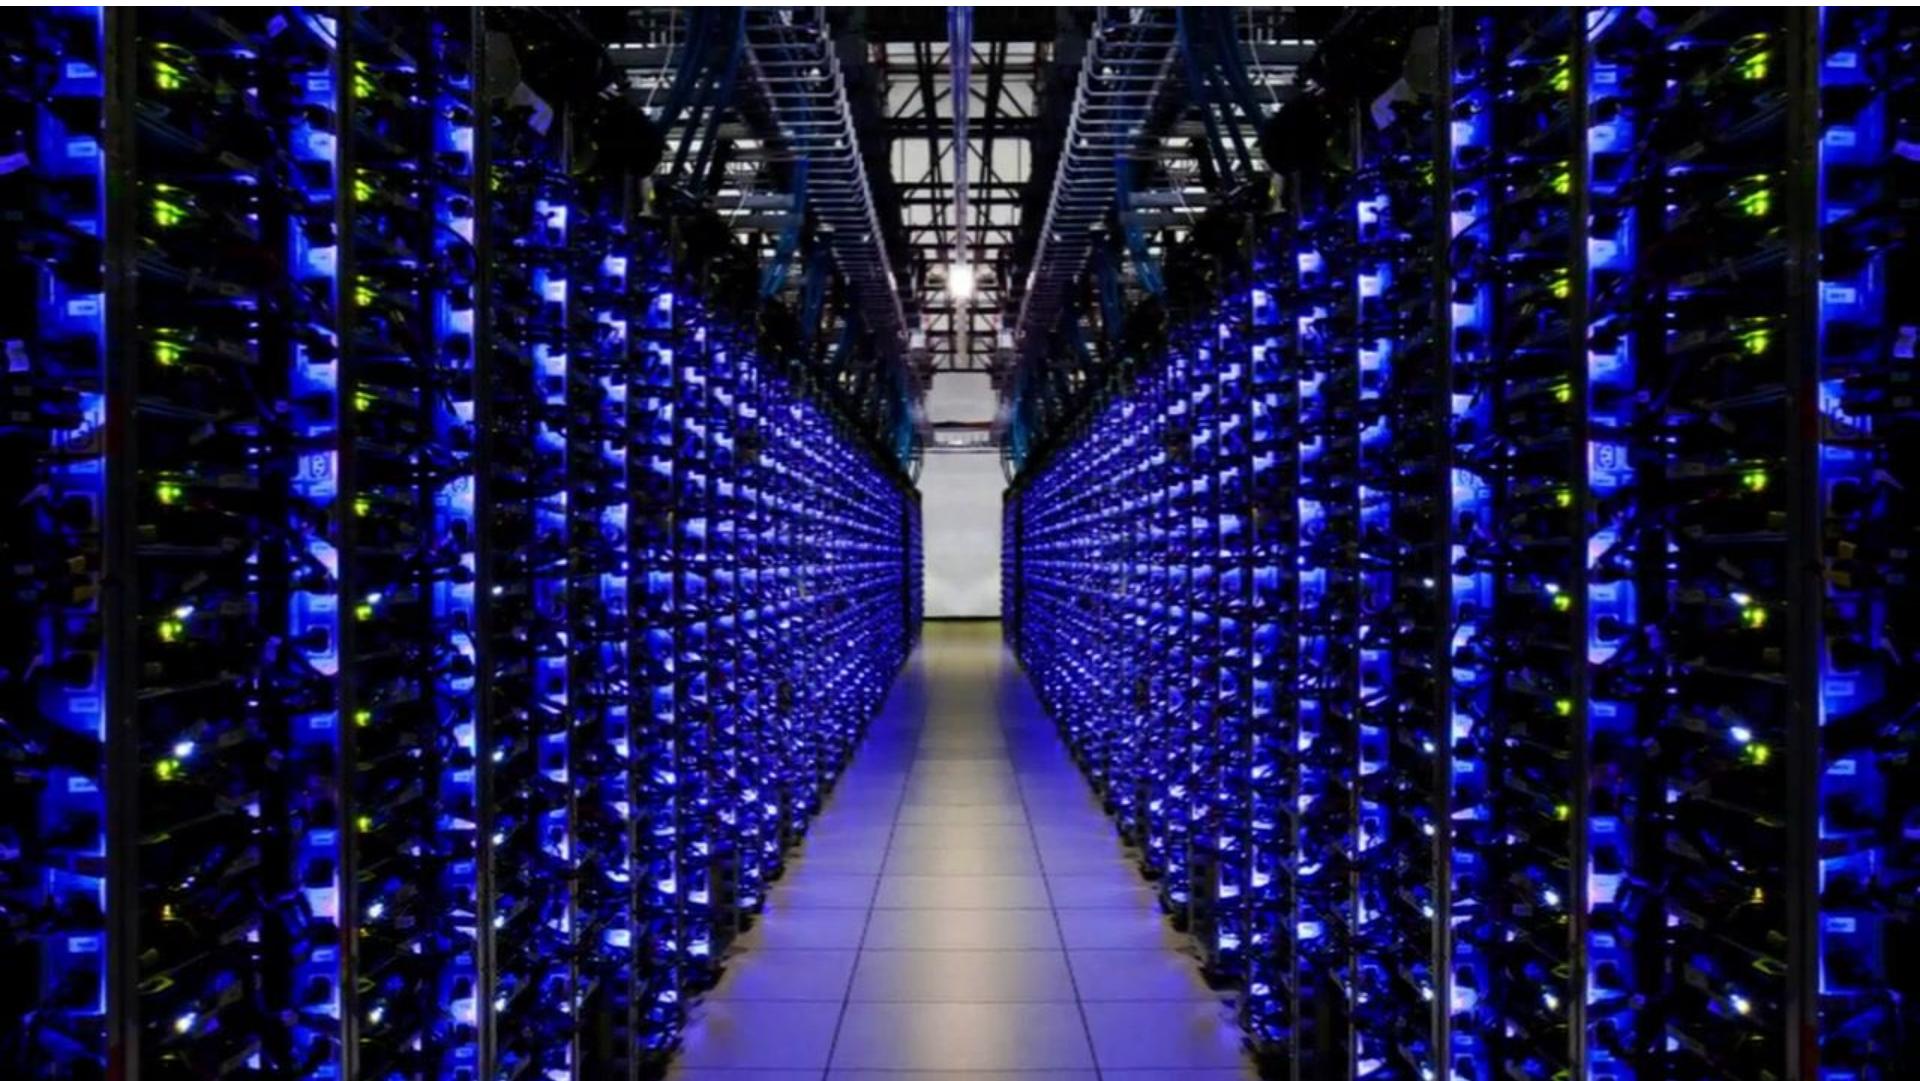
99% -- same as humans

Some cases that it gets wrong



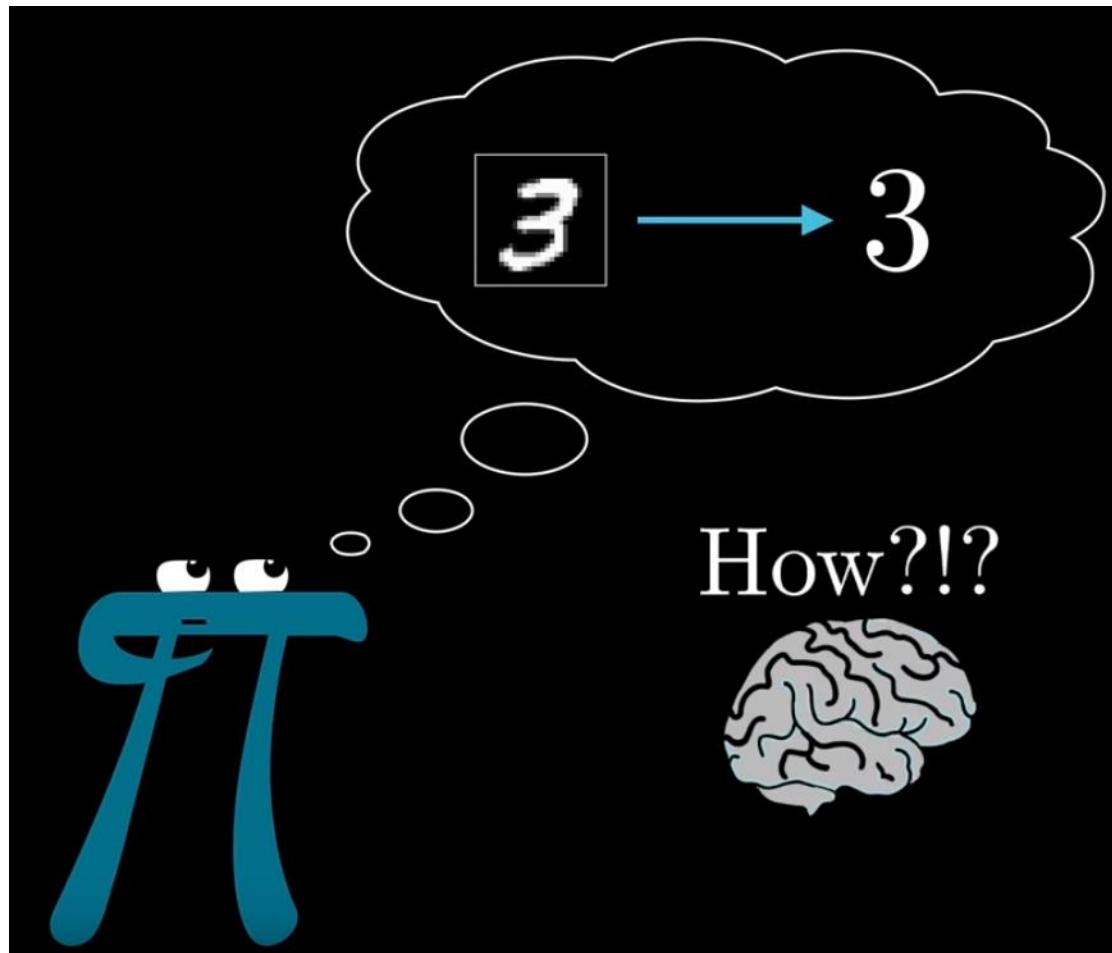
correct classification is in the top right;
program's classification is in the bottom right

Lots of computational power needed



Another informal intro to NN

<https://www.youtube.com/watch?v=aircAruvnKk>



Success Stories of AI

- Play professionally chess and many other games.
- Discover and prove theorems
- Drive a car
- Do image recognition (with accuracy often better than humans)
- Do speech recognition
- Do automatic translation

That's all so far -- questions?

